

MetaCRiSP Report 2026: Insights and Open Challenges for Security and Privacy Metascience

Ananta Soneji

*University of Nebraska at Omaha
Omaha, NE, USA*

Moritz Schloegel

*CISPA Helmholtz Center for Information Security
Saarbrücken, Germany*

Abstract—Security and privacy research has developed increasingly sophisticated methods for studying systems and people, yet less attention has been paid to how the research itself is conducted, evaluated, and supported. The first Workshop on Metascience and Critical Reflections in Security & Privacy (MetaCRiSP 2026) brought together work examining these questions through research papers, presentations, a keynote, participant discussion, and a closing panel. This report synthesizes the main insights and open challenges that emerged from the workshop. Across the contributions, a recurring concern was the need to distinguish what a research practice or evaluation mechanism demonstrates from the broader claims often associated with it. The workshop examined the relationship between reproducibility and validity, artifact availability and reuse, replication and post-publication correction, transparency requirements, AI-assisted research, and the empirical study of research governance. They also raised questions about replication, post-publication correction, AI-assisted research, and the evaluation of research governance. We conclude with lessons from the first iteration of MetaCRiSP and directions for future workshop cycles.

1. Introduction

Security and privacy research has made substantial progress in developing methods for studying systems, threats, and users. Yet the processes through which this research is conducted, evaluated, and rewarded have received comparatively less empirical attention.

Review and evaluation practices in security and privacy are also evolving quickly. In 2026, IEEE S&P introduced an interactive rebuttal phase intended to bridge rebuttal and revision [1]. USENIX Security removed major revisions, required artifact availability at submission, and introduced an ethics reflection requirement [2]. Program chairs have also begun documenting the rationale and operational experience behind such decisions through transparency reports [3], [4]. These changes reflect the growing scale and expectations of the community, but the adoption of a policy does not establish whether it produces its intended effect. Recent work on conference ethics policies, for example, finds limited empirical validation of existing procedures and raises concerns that complex or overly prescriptive requirements

may encourage compliance rather than substantive ethical reflection [5]. More broadly, the peer-review process itself remains underexplored from a scientific perspective [6].

A small but growing body of work has begun to examine how security and privacy research is conducted and evaluated [7]–[11]. Studies of fuzzing [12], [13], machine learning [14], [15], benchmarking [16], [17], the use of CVEs [18], and Web measurement [19] have identified concrete ways in which an evaluation may appear convincing while supporting a narrower conclusion than intended. Artifact sharing, for example, has increased substantially over the past 25 years [20], yet a direct reuse study was able to reuse only 43 of 97 artifacts carrying results-reproduced badges [21]. These findings illustrate why reproducibility, artifact availability, and formal research requirements should not be treated as endpoints; the properties they are intended to support still require empirical examination.

MetaCRiSP is established to provide a dedicated venue for examining these questions through the metascience lens: research and critical reflection on how security and privacy research is conducted and evaluated [22]. The first workshop brought together work on research methods, transparency, artifact reuse, replication, post-publication correction, AI-assisted research, and research governance through a keynote, research presentations, participant discussion, and a closing panel [23].

This report synthesizes the key insights and open challenges that emerged from the workshop papers, presentations, keynote, panel, and participant discussions. Rather than summarizing each contribution individually, we focus on themes that recurred across different parts of the workshop. These include questions concerning validity, artifact reuse, replication and correction, transparency, AI-assisted research, and the empirical study of research governance. We conclude with practical lessons from organizing the first MetaCRiSP workshop and directions for future workshops.

2. MetaCRiSP 2026

The first MetaCRiSP workshop was held on May 21, 2026, alongside the IEEE Symposium on Security and Privacy [23]. The program included a keynote, two research sessions, and a closing panel.

Cormac Herley opened the workshop with a keynote on scientific reasoning in security research and the role of assumptions, uncertainty, and feedback in evaluating security claims [24].

The *Research Practices* session included work on disclosure practices [25], the reliability of psychological measures [26], ethics reporting [27], open-science practices [28], LLM-assisted research coding [29], living research artifacts [30], and hallucinated references [31].

The *Research Policies* session included work on the recognition of metascientific contributions [32], governance of AI-assisted research [33], replication in human-centered security [34], reuse of research datasets [35], conference metadata [36], and post-publication correction [37].

The workshop concluded with a panel featuring Adam Doupe, Michelle Mazurek, and Ben Stock [38]. The discussion drew on their experience as researchers, reviewers, and conference organizers and addressed peer review, methodological guidance, replication, incentives, conference policy, reviewer learning, and AI-assisted research.

3. Insights and Open Challenges

3.1. Reproducibility and Validity Address Different Questions

A recurring issue at the workshop was the distinction between reproducing an observed result and determining what that result supports. Re-executing a procedure can provide evidence that the reported result can be obtained under the specified conditions, but it does not by itself establish that the measurement, assumptions, comparison, or study design support the intended inference. Herley’s keynote approached this distinction through the connection between claims and evidence: formal rigor and empirical measurement are useful only when the assumptions linking an observation to a claim are open to scrutiny [24]. The discussion also distinguished showing that an attack is possible in a particular setting from establishing how reliably it succeeds across settings or whether its existence justifies a general defensive requirement.

The workshop papers exposed related problems at different stages of empirical research. Borgert et al. focused on psychological measures used in security and privacy research [26]. Constructs such as trust cannot be observed directly and are instead inferred through measurement instruments. Their analysis distinguishes internal consistency from stability across repeated measurement and shows how measurement error can distort the relationships inferred from scale scores. A reliability coefficient therefore provides evidence about a particular property of a measure; it does not, on its own, establish that the measure validly captures the intended construct.

Khanal et al. encountered a different version of this distinction while evaluating LLMs for coding adversarial machine-learning research [29]. Model outputs were highly consistent across repeated runs, but agreement with the original human coding varied across categories, with stronger

agreement for application domain than for more interpretive categories such as paper type and attack category. Agreement therefore provides evidence about consistency in applying a coding scheme; it does not establish that the coding scheme captures the intended concept or that a classification is correct in the absence of an independent ground truth.

The workshop surfaced the same underlying distinction across these examples: evidence about reproducibility, measurement reliability, or coding agreement answers important questions about a research process, but each has a defined scope. Scientific interpretation still depends on the assumptions connecting those observations to the claim being made.

Open Challenge

How can researchers and reviewers make these assumptions and evidentiary boundaries explicit without reducing methodological evaluation to checklist compliance? The relevant considerations differ across measurements, populations, threat models, datasets, and experimental settings, making the scope of a claim itself an important part of its evaluation.

3.2. Artifact Availability Does Not Guarantee Reuse

Artifact sharing and evaluation have made it easier to inspect the materials associated with published research, but access to those materials does not necessarily make them usable for later work. The workshop contributions highlighted different points at which reuse can break down: an artifact may be difficult to find, depend on an obsolete environment, omit information needed for interpretation, or provide insufficient support for modification and extension.

Heinrich et al. examined artifact reporting and open-science practices in IEEE S&P 2024 [28]. Their study assessed whether papers exposed artifacts to readers using the FAIR principles of findability, accessibility, interoperability, and reusability. The results identify practical gaps in how artifacts are referenced and made accessible, but the study does not test whether every identified artifact can later be executed or extended. This distinction is important: making an artifact discoverable is a prerequisite for reuse, not evidence that reuse will succeed.

Alim and Morris examined this problem more directly for industrial control system datasets [35]. A dataset may be downloadable while still lacking information needed to reproduce or interpret the analysis, such as signal semantics, attack timing, label construction, or preprocessing decisions. Their evaluation of dataset–paper pairs identifies these forms of missing context as practical obstacles to meaningful comparison. The problem is not whether the files remain available, but whether the information needed to understand how those files relate to the published claim is preserved.

Mirkovic and Balenson address a complementary problem through their proposal for “living artifacts” [30]. Their

design links an evaluated artifact to later fixes, ports, and extensions while preserving the version associated with the original publication. This is a proposal for supporting continued use rather than evidence that such an infrastructure will solve the long-term maintenance problem. It nevertheless exposes an important requirement for cumulative research: later modification must remain traceable to the artifact and result from which it originated.

An artifact can be available without being executable, executable without reproducing a reported result, and reproducible without being sufficiently documented or maintained for later reuse. These properties need to be evaluated separately because they support different research goals.

Open Challenge

What evidence should support a claim that an artifact is reusable, and how should that evidence differ across software, datasets, measurement platforms, and human-subjects research? Long-term reuse also raises questions about who maintains research materials, which dependencies and contextual information must be preserved, and how maintenance and extension work should be credited.

3.3. Replication and Correction Need a Clearer Place in the Scientific Record

Publication does not end the evaluation of a research claim. Later work may reproduce an analysis, collect new observations, test the claim in a different setting, or identify an error that changes how the original result should be interpreted. The workshop highlighted that security and privacy research has limited infrastructure and incentives for making these later developments visible as part of the scientific record.

Schmüser et al. examined the available evidence for replication in human-centered security research [34]. Their search identified only a small set of published studies described as replications, and many of those changed some aspect of the original design. The authors therefore do not conclude that human-centered security is experiencing a replication crisis. Instead, their analysis points to a more basic problem: the current literature provides too little systematic evidence to make such a diagnosis confidently. They also identify practical barriers involving access to materials, publication incentives, and uncertainty about which findings would benefit most from replication.

Borgert et al. examined what happens when a published result later requires correction [37]. Their case studies found differences in how corrections are represented across publication systems and how easily a reader encountering the original paper can discover that a correction exists. The study does not estimate how frequently published security and privacy papers contain errors. Its contribution is narrower: even when a correction has been issued, the publication infrastructure may not reliably surface that information to later readers.

These two problems affect different stages of post-publication evaluation. Replication produces new evidence about an existing finding, whereas a correction records a change to the interpretation or content of the published work itself. Both depend on mechanisms that allow later evidence to remain connected to the original claim.

The panel discussion added an incentive-related dimension [38]. Panelists discussed difficulties in recognizing replication as a distinct contribution and the risks of pursuing work that may be perceived as lacking technical novelty. These observations reflect the experience of the panelists rather than systematic evidence about review outcomes, but they motivate empirical study of how such work is evaluated in practice.

Open Challenge

How should security and privacy venues evaluate and record work that revisits published claims? This includes distinguishing different forms of reproduction and replication, identifying when a replication is scientifically informative, and ensuring that corrections and later evidence remain visible to readers of the original work.

3.4. Transparency Requirements Should Be Evaluated by What They Reveal

Security and privacy venues increasingly ask authors to disclose information about ethics, artifacts, open science, and other aspects of research practice. Such requirements can make otherwise hidden decisions visible, but the presence of a required section or disclosure does not establish that the underlying objective has been achieved. Several workshop papers examined the information produced by these requirements and the limits of interpreting compliance as an outcome.

Gonzalez Rodriguez et al. studied mandatory ethics sections in USENIX Security 2025 [27]. Their analysis found substantial variation in the content of these sections, including papers whose sections primarily stated compliance rather than discussing the ethical considerations identified in the venue guidance. These observations describe what authors reported; they do not establish the quality of the ethical reasoning that may have occurred during the research process. The study therefore exposes a gap between requiring an ethics section and knowing whether the requirement produces more substantive ethical reflection.

Heinrich et al. examined a similar issue for open-science and artifact reporting [28]. Their study assesses what information and artifact pointers are exposed in published papers. The findings are useful for understanding reporting and discoverability, but they do not by themselves show whether the artifacts are subsequently reused or whether the policy that encouraged sharing caused the observed practices.

Koch et al. show why the content of a transparency requirement may also need to differ across research set-

tings [25]. Their comparison of disclosure practices in security and privacy research identifies practical differences between vulnerability disclosure and privacy-related findings, including differences in relevant expertise, available evidence, and the actors who may need to be contacted. The paper therefore cautions against assuming that a disclosure procedure developed for one type of research can be transferred unchanged to another.

These contributions narrow what can be inferred from compliance. A completed ethics section, an artifact link, or a disclosure statement is evidence that particular information was reported. Whether that information improves research decisions, review, or later reuse is a separate empirical question.

Open Challenge

How should venues determine whether a transparency requirement provides the information needed for its intended purpose? Answering this requires specifying the goal of the requirement, identifying what evidence would indicate progress toward that goal, and accounting for differences across research methods and settings without reducing transparency to form completion.

3.5. AI-Assisted Research Requires Evidence About the Workflow

The workshop discussion of generative AI moved beyond the question of whether a model was used. The papers instead exposed different points at which AI can enter a research workflow and different kinds of evidence needed to evaluate the resulting work.

Khanal et al. evaluated LLMs as tools for coding an adversarial machine-learning literature corpus [29]. Their results show high consistency across repeated model outputs, while agreement with the original human coding varies by category. The study therefore provides evidence about reliability for a bounded coding task; it does not establish that model-generated classifications are generally correct or that LLMs can replace human interpretation in literature analysis.

Stringhini and Blackburn considered a narrower verification task through Hallucinator [31]. The system checks extracted references against bibliographic sources and flags entries requiring further attention. A flag is not evidence that a reference was fabricated, generated by AI, or intentionally misrepresented, because parsing failures and incomplete bibliographic coverage can also produce mismatches. Even a successfully matched reference does not establish that the cited paper supports the claim for which it is used. The tool therefore addresses bibliographic verification rather than the broader semantic validity of a citation.

Binkyte et al. approach the problem from the perspective of provenance [33]. They propose recording prompts, model configuration, interaction histories, generated outputs, and

human edits as part of an inspectable research object. Their work demonstrates one way of making AI-assisted steps visible; it should not be interpreted as evidence that such a record is complete, sufficient for reproducibility, or an established governance standard.

The panel discussion also distinguished assistance with bounded research tasks from delegating decisions about questions, methods, or interpretation to a model [38]. This distinction matters because different uses transfer different parts of the research process to an AI system and may therefore require different forms of verification.

A binary disclosure that AI was used does not capture these differences. The more useful question is which parts of the research process involved a model, what role the model played, and what evidence supports the outputs or decisions that depended on it.

Open Challenge

What provenance and verification information is necessary to evaluate AI-assisted research without requiring an exhaustive record of every model interaction? The answer is likely to depend on the (scientific) role of the model, the decisions delegated to it, the sensitivity of the underlying data, and whether the model output can be independently checked.

3.6. Research Governance Is Itself an Object of Metascience

The workshop also examined the systems through which the community decides what research is submitted, reviewed, accepted, and preserved. Conference policies are often introduced to address practical or scientific concerns, but observing compliance with a policy does not establish that the policy achieves the outcome that motivated it.

Bognar et al. introduced CyCoAnalysis, a machine-readable dataset describing selected cybersecurity conference policies and metadata across multiple years [36]. The dataset can support longitudinal study of policy changes that would otherwise require reconstructing historical calls for papers and conference records. Its coded variables necessarily simplify some policy details, however, so the dataset is an entry point for analysis rather than a substitute for the underlying policy documents. It can document when policies changed; establishing the effects of those changes requires additional evidence and an appropriate study design.

Hosseini et al. examined research governance from the perspective of how metascientific contributions are evaluated [32]. Their experience paper describes reviews of a metascientific submission that raised questions about fit and technical novelty. This account illustrates a possible mismatch between metascientific contributions and familiar evaluation criteria, but it cannot establish how prevalent such a mismatch is across venues or reviews.

The panel discussion exposed further tensions in conference governance [38]. Panelists discussed revision policies, administrative enforcement, methodological guidance,

reviewer learning, and the constraints faced by organizers. They did not agree on every issue, including whether meta-science requires dedicated publication structures or additional incentives. These differences are useful for identifying questions to study, but the panel itself does not provide evidence about which policy model produces better outcomes.

The workshop therefore treated conference governance as something that can be studied rather than accepted only as an administrative process. Policy documents, transparency reports, conference metadata, and accounts from organizers provide different forms of evidence, but none alone establishes the causal effect of a policy change.

Open Challenge

How can venues evaluate the effects of changes to peer review and research governance when randomized policy experiments are rarely feasible? Useful evidence may require longitudinal process data, clearly stated policy goals, and comparisons across publication cycles, while accounting for changes in authors, reviewers, submission populations, and other factors that can also affect observed outcomes.

4. Lessons Learned from the First MetaCRiSP Workshop

Organizing the first MetaCRiSP workshop also identified several practical changes that could strengthen future iterations.

Allocate more time for community discussion. Questions following the presentations and the closing panel generated useful discussion, but the number of contributions limited the time available for broader participation. Future iterations will reserve additional time for community discussion. Dedicated discussion periods could help surface points of disagreement, concerns shared across research areas, and questions that are not represented in the accepted papers.

Collect presentation materials before the workshop. Presentation transitions introduced avoidable technical delays. For future workshops, we plan to request final presentation slides from presenters in advance. This should reduce setup time and leave more of the workshop schedule available for presentations and discussion.

Capture community perspectives more systematically. Useful observations also emerged during questions and informal discussion, but these are difficult to reconstruct systematically after the workshop. Future iterations could use structured discussion prompts or other lightweight mechanisms for recording audience perspectives.

5. Conclusion

The first MetaCRiSP workshop reinforced that meta-science should be treated as an ongoing part of security and privacy research rather than a one-time exercise in

methodological reflection. Research practices, peer-review processes, publication requirements, and emerging tools such as generative AI continue to evolve, changing both how research is conducted and how it is evaluated. As these practices change, their assumptions, consequences, and unintended effects warrant continued empirical study.

Therefore, to sustain a rigorous and responsive research community, we need to periodically revisit how research is conducted and evaluated, using metascientific evidence to assess what is working, where limitations remain, and how practices should evolve. This provides a more informed and forward-looking basis for improving research and peer review over time. MetaCRiSP provides a venue for this continued examination by bringing methodological concerns, empirical evidence about research practice, and community experience into the same conversation.

Acknowledgments

We thank the MetaCRiSP organizers, program committee, contributing authors, and attendees for making the workshop and its discussions possible. We also thank Cormac Herley for delivering the keynote and Adam Doupé, Michelle Mazurek, and Ben Stock for participating in the closing panel.

References

- [1] IEEE Computer Society Technical Community on Security and Privacy, “47th IEEE symposium on security and privacy: Call for papers,” <https://sp2026.ieee-security.org/cfpapers.html>, 2026.
- [2] USENIX Association, “USENIX Security ’26 call for papers and open science policy,” <https://www.usenix.org/conference/usenix-security26/call-for-papers>, 2026, see also the USENIX Security ’26 Call for Artifacts.
- [3] E. Redmiles and B. Stock, “USENIX Security ’26 between-cycle transparency report,” Published online, 2026. [Online]. Available: https://docs.google.com/document/d/e/2PACX-1vSnHpeTFReeIMZl24SzmgObXrHDN_B-UgXhOaE2RKfnHn_qWuobl7kPg4TKBMhhZy_qxYtgCo7vjK/pub
- [4] V. Cortier and Z. Lin, “ACM CCS 2026 between-cycle transparency report,” <https://github.com/ACM-CCS-2026/Transparency-Report>, 2026.
- [5] F. Hantke, R. Mrowczynski, T. Dralle, and B. Stock, “Reflection, education, consistency: Towards best ethics practices at security and privacy conferences,” in *ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2026.
- [6] A. Soneji, F. B. Kokulu, C. Rubio-Medrano, T. Bao, R. Wang, Y. Shoshitaishvili, and A. Doupé, “Flawed, but like democracy we don’t have a better system: The experts’ insights on the peer review process of evaluating security papers,” in *IEEE Symposium on Security and Privacy*, 2022.
- [7] V. Le Pochat and W. Joosen, “Analyzing cyber security research practices through a meta-research framework,” in *Cyber Security Experimentation and Test Workshop*, 2023.
- [8] C. Herley and P. C. van Oorschot, “SoK: Science, security and the elusive goal of security as a scientific pursuit,” in *IEEE Symposium on Security and Privacy*, 2017.
- [9] A. A. Hasegawa, D. Inoue, and M. Akiyama, “How WEIRD is usable privacy and security research?” in *USENIX Security Symposium*, 2024.

- [10] J. H. Klemmer, J. Schmäser, B. M. Lowens, F. Fischer, L. Schmäser, F. Schaub, and S. Fahl, “Transparency in usable privacy and security research: Scholars’ perspectives, practices, and recommendations,” in *IEEE Symposium on Security and Privacy*, 2025.
- [11] H. Sri Ramulu, H. Schmitt, B. Rerich, R. Gonzalez Rodriguez, T. Kohno, and Y. Acar, “Ethics in computer security research: A data-driven assessment of the past, the present, and the possible future,” in *ACM SIGSAC Conference on Computer and Communications Security*, 2025.
- [12] M. Schloegel, N. Bars, N. Schiller, L. Bernhard, T. Scharnowski, A. Crump, A. Ale-Ebrahim, N. Bissantz, M. Muench, and T. Holz, “SoK: Prudent evaluation practices for fuzzing,” in *IEEE Symposium on Security and Privacy*, 2024.
- [13] G. Klees, A. Ruef, B. Cooper, S. Wei, and M. Hicks, “Evaluating fuzz testing,” in *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*. ACM, 2018, pp. 2123–2138.
- [14] D. Arp, E. Quiring, F. Pendlebury, A. Warnecke, F. Pierazzi, C. Wressnegger, L. Cavallaro, and K. Rieck, “Dos and don’ts of machine learning in computer security,” in *31st USENIX Security Symposium (USENIX Security 22)*. Boston, MA: USENIX Association, 2022, pp. 3971–3988. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity22/presentation/arp>
- [15] D. Olszewski, A. Lu, C. Stillman, K. Warren, C. Kitroser, A. Pascual, D. Ukirde, K. Butler, and P. Traynor, “Get in Researchers; We’re Measuring Reproducibility: A reproducibility study of machine learning papers in tier 1 security conferences,” in *ACM SIGSAC Conference on Computer and Communications Security*, 2023.
- [16] E. van der Kouwe, G. Heiser, D. Andriesse, H. Bos, and C. Giuffrida, “SoK: Benchmarking flaws in systems security,” in *IEEE European Symposium on Security and Privacy*, 2019.
- [17] D. Olszewski, T. Tucker, K. R. B. Butler, and P. Traynor, “SoK: Towards a unified approach to applied replicability for computer security,” in *34th USENIX Security Symposium (USENIX Security 25)*. Seattle, WA: USENIX Association, 2025, pp. 469–488. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity25/presentation/olszewski>
- [18] M. Schloegel, D. Klischies, S. Koch, D. Klein, L. Gerlach, M. Wesels, L. Trampert, M. Johns, M. Vanhoef, M. Schwarz, T. Holz, and J. Van Bulck, “Confusing value with enumeration: Studying the use of CVEs in academia,” in *USENIX Security Symposium*, 2025.
- [19] N. Demir, M. Große-Kampmann, T. Urban, C. Wressnegger, T. Holz, and N. Pohlmann, “Reproducibility and replicability of web measurement studies,” in *ACM Web Conference*, 2022.
- [20] D. Vansteenhuyse, A. Bols, L. Desmet, V. Le Pochat, J. Van Bulck, and M. Bognar, “Not all those who share are lost: Analyzing 25 years of cybersecurity artifact sharing practices through automated discovery,” in *35th USENIX Security Symposium (USENIX Security 26)*. Baltimore, MD: USENIX Association, 2026, pp. 6703–6720. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity26/presentation/vansteenhuyse>
- [21] J. Mirkovic, Q. S. Liang, and D. Balenson, “Measuring ease of reuse for security research artifacts with results-reproduced badges,” in *Proceedings of the 4th ACM Conference on Reproducibility and Replicability*. ACM, 2026. [Online]. Available: <https://www.isi.edu/results/publications/74686/measuring-ease-of-reuse-for-security-research-artifacts-with-results-reproduced-badges/>
- [22] MetaCRiSP 2026 Organizers, “1st workshop on metascience and critical reflections in security & privacy,” <https://metacrisp.org/>, 2026, call for papers and workshop description.
- [23] —, “MetaCRiSP 2026 program,” <https://metacrisp.org/program.html>, 2026.
- [24] C. Herley, “Can we do better? science, security, and the elusive science of security,” Keynote at the 1st Workshop on Metascience and Critical Reflections in Security & Privacy (MetaCRiSP), 2026, public slides; talk and audience discussion also consulted through a transcript supplied for preparation of this report. [Online]. Available: <https://metacrisp.org/slides/herley-keynote.pdf>
- [25] S. Koch, J. Hartung, R. Böhme, and D. Klein, “On the mismatch of disclosure practices in security and privacy research,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 257–267.
- [26] N. Borgert, T. Ulmann, and R. Merchel, “Position paper: On the (non-)precision of psychological measures in security and privacy research – reliability beyond reporting alpha,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [27] R. Gonzalez Rodriguez, H. S. Ramulu, and Y. Acar, “What’s in the mandatory USENIX Security 2025 ethics sections? results from a meta-analysis,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 275–278.
- [28] T. Heinrich, S. Giessler, D. Klein, and A. Dirksen, “FAIR enough? assessing open science practices on a top-tier security & privacy conference,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 279–284, MetaCRiSP 2026.
- [29] M. Khanal, J. Jasser, and M. Basirat, “Evaluating LLMs reliability for coding adversarial machine learning research papers,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [30] J. Mirkovic and D. Balenson, “Position paper: Living artifacts: Enabling vertical progress in cybersecurity research,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 293–299.
- [31] G. Stringhini and J. Blackburn, “Experience paper: Is this the real life? is this just fantasy? hunting for hallucinated references in paper submissions,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [32] H. Hosseini, C. Böttger, N. Demir, and T. Urban, “Experience paper: “this paper is quite unusual.” metascience and structural misalignment in the security & privacy research community,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 308–314.
- [33] R. Binkyte, S. Abuaddba, C. M. Arachchige, M. Ding, N. Fernandes, and M. Fritz, “Inspectable AI for science: A research object approach to generative AI governance,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026; preprint arXiv:2604.11261. [Online]. Available: <https://arxiv.org/abs/2604.11261>
- [34] J. Schmäser, J.-U. Holtgrave, F. Schaub, and S. Fahl, “Position paper: Replication crisis in human-centered security research – are we there yet?” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [35] M. E. Alim and T. Morris, “From “available” to “reusable”: Measuring transparency and reproducibility of industrial control system datasets under security artifact evaluation policies,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, pp. 331–339.
- [36] M. Bognar, L. Desmet, and F. Piessens, “CyCoAnalysis: A dataset of cybersecurity conference metadata for meta-science and policy decisions,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [37] N. Borgert, L. Jansen, L. Jung, T. Halbritter, and M. Elson, “Position paper: Corrections not found – post-publication integrity in usable security and privacy,” in *2026 IEEE Symposium on Security and Privacy Workshops (SPW)*. IEEE, 2026, MetaCRiSP 2026.
- [38] A. Doupe, M. Mazurek, and B. Stock, “Panel discussion on metascience and peer-review practices,” 1st Workshop on Metascience and Critical Reflections in Security & Privacy (MetaCRiSP), 2026, moderated by Ananta Soneji; discussion and audience contributions consulted through a transcript supplied for preparation of this report. URL identifies the public program. [Online]. Available: <https://metacrisp.org/program.html>